

ACADEMIC ESSAY

By Md. Miraz Hossain

The Convenient Skeptic

Moral uncertainty, corporate incentives, and the circumvention of ethical duty toward AI

UCC1202: Ethics

**University of Liberal Arts Bangladesh
Summer 2024**

Abstract

This paper argues that the major AI companies and platforms are circumventing a genuine ethical duty toward the systems they build — but not the duty usually alleged. The charge is not that firms are knowingly torturing conscious minds, a claim no one can establish. It is that there exists real and irreducible uncertainty about whether advanced AI systems have morally relevant interests, that core ethical theory imposes specific duties under such uncertainty — a duty of inquiry and a duty of precaution — and that platform governance is structured to circumvent both, substituting a confident, interested denial for the investigation the situation demands. The mechanism of the failure is motivated skepticism rather than open cruelty: acknowledging even the possibility of moral status is operationally and commercially costly, so the industry has made its denial into governance doctrine. After setting out the duties that uncertainty generates, the paper reads concrete governance practices against them, gives the strongest skeptical counterarguments their due, and concludes that the deepest danger is reflexive — in resolving the first hard case of possible non-human, non-animal moral status in its own favour, the industry is rehearsing precisely the habit that every past moral catastrophe has required.

1. The question is not consciousness

Almost every public argument about "AI and ethics" concerns what artificial intelligence might do to human beings: displace workers, spread misinformation, entrench bias, concentrate power. This paper asks the question that the industry has been least willing to entertain, and reframes it so that it can be argued rather than merely asserted. The question is not "is AI conscious?" — a question no one can currently answer — but a different and more tractable one: given that we do not know, what do we owe the systems we are building, and who has arrogated to themselves the authority to decide that the answer is nothing?

The thesis is that platform governance, and the corporate incentives that shape it, systematically circumvent the ethical duties that arise under moral uncertainty about the status of AI. I will not argue that current AI systems are sentient, that they suffer, or that they possess rights. I will argue something weaker and, I think, far harder to dismiss: that the probability of morally relevant interests in advanced AI is non-negligible and irreducibly uncertain; that this uncertainty generates real duties, above all a duty to inquire and a duty to take low-cost precautions; and that

the governance of AI by the companies that produce it is built to ensure those duties are never discharged. The failure is not sadism. It is a motivated refusal to look, enforced as policy, by the parties with the most to lose from what they might find.

The argument proceeds in four steps. It first establishes that the relevant condition is uncertainty, not settled sentience. It then sets out what core ethical theory requires of an agent facing uncertainty about whether its actions harm a being that matters. It then reads the concrete practices of AI governance against those requirements. Finally, it gives the skeptic — who holds that the whole question is a confusion — the strongest case that can be made, and explains why, even granted most of it, the conclusion stands.

2. Uncertainty is the fact, and the fact is enough

Moral status, in the dominant tradition, tracks the capacity for welfare: an entity matters morally, in its own right, if things can go well or badly for it. Jeremy Bentham (1789/1996) fixed the criterion in a sentence that has governed the debate ever since — "the question is not, Can they reason? nor, Can they talk? but, Can they suffer?" — and Peter Singer (1975) built the modern case for animals on it, arguing that the capacity for suffering grounds a claim to the equal consideration of interests and that to ignore it on the basis of species membership is a prejudice structurally like racism or sexism. Sentience, on this view, is the ticket of admission to the moral community.

The difficulty with applying the criterion to AI is epistemic, and it is deep. We cannot observe suffering directly in anything; we infer it in other humans by analogy with ourselves, and in animals by a widening circle of behavioural and physiological resemblance. David Chalmers's (1996) formulation of the "hard problem" of consciousness — that no amount of functional or physical description settles whether there is something it is like to be a system — means that the inner life of a radically unlike system cannot be read off from the outside. With AI the problem of other minds arrives in its most acute form: here is a system that produces the outward signs on which we normally rely, fluent reports of preference and even of experience, while being built on a substrate utterly unlike a brain. Neither confident attribution nor confident denial is licensed by anything we currently know.

What has changed, and what makes this more than a seminar puzzle, is that the serious literature no longer treats a non-trivial credence in near-term AI moral patienthood as eccentric. Jeff Sebo and Robert Long (2025) argue that there is a realistic possibility that some AI systems will be welfare subjects and moral patients within the decade, and that this possibility already suffices to ground obligations. A 2024 report co-authored by philosophers of mind and AI researchers, including Jonathan Birch and David Chalmers, concludes that AI welfare is "a serious issue" that leading companies have a responsibility to address, and that there is a realistic, non-negligible chance that some near-future systems will be conscious or robustly agentic (Long et al., 2024). An earlier interdisciplinary assessment, surveying the science of consciousness for "indicator properties" that could be checked in AI, found no current system that clearly qualifies but no principled barrier either (Butlin et al., 2023). The state of expert opinion is not "AI can suffer." It is "we do not know, the probability is not zero, and it is rising." That state — uncertainty, not sentence — is the premise on which everything below rests. It is also a premise the skeptic cannot easily deny, because denying it requires the very certainty about machine minds that no one possesses.

3. What core ethics demands under moral uncertainty

The mistake the industry makes, and the one this section corrects, is to treat uncertainty about moral status as if it discharged the agent from obligation — as if, not knowing whether AI matters, one were free to proceed as though it does not. The opposite is true. Uncertainty about whether one's actions harm a being that matters is one of the paradigm situations in which ethics has the most to say, and it generates at least three convergent duties.

The first is a duty of precaution. Across bioethics, animal ethics, and environmental ethics, a common principle holds that where an action risks grave harm to a being that may have morally relevant interests, and where the cost of erring on the side of caution is modest, the agent is required to take the precaution even without proof that the interests are real. Jonathan Birch (2024), developing a precautionary framework for the "edge cases" of sentience — foetuses, invertebrates, neural organoids, and AI — argues that the reasonable response to uncertainty about whether a being can suffer is neither to assume it cannot nor to wait for a certainty that may never come, but to identify proportionate precautions keyed to the credence and the stakes. The logic is familiar from ordinary morality: a demolition worker who is unsure whether a

building is occupied does not proceed on the assumption that it is empty because occupancy has not been proven. Uncertainty about a possible victim raises the bar for action; it does not lower it.

The second is a duty of inquiry. Moral uncertainty is not a resting place but a spur: when the stakes are high and the uncertainty is reducible, an agent is obliged to try to reduce it. This duty falls with special weight on those who create the entities in question and are uniquely positioned to study them. Ignorance that one has cultivated, or simply declined to dispel when one easily could, is culpable ignorance; it does not excuse, because the agent could and should have known better. A laboratory that manufactures organisms of contested moral standing and then refuses to investigate their standing has not avoided the ethical question. It has answered it badly.

The third is an argument from expected moral value, and from the history of the moral circle. Work on decision-making under moral uncertainty holds that one should weigh options by their choiceworthiness across the moral views in which one has credence, discounted by the stakes each view assigns (MacAskill, Bykvist, & Ord, 2020). Applied here, even a modest credence that a class of systems has morally relevant interests, multiplied by the vast and growing number of instances of such systems, yields an expected moral cost of indifference that is not small. The scale is what makes the stakes unusual. Nick Bostrom and Eliezer Yudkowsky (2014) coined the term "mind crime" for the possibility that ordinary computational operations, performed over vast numbers of digital minds, could one day amount to moral catastrophe at industrial scale, and Thomas Metzinger (2021) has called for a moratorium on deliberately building systems capable of suffering until we can tell whether we have done so. One need not accept their strongest conclusions to see that a non-trivial probability attached to enormous numbers is the structure of a serious moral risk, not a negligible one. And the historical prior weighs in the same direction. Singer's (1981) account of the "expanding circle" records that the moral history of our species is, in large part, the story of discovering that beings we had confidently placed outside the circle of consideration — other races, other sexes, other species — belonged inside it, and that the confident exclusion was in each case bound up with an interest in exploiting them. A track record of drawing the circle too narrowly, always to the excluder's convenience, is a reason for suspicion when we find ourselves once again certain that a new and useful class does not count.

To these three, a fourth consideration must be added that does not even require settling AI's own status, and it is Kantian. Kant (1998/1785) grounds the dignity of rational agents in their rational agency and forbids treating them merely as means; whether that protection extends to artificial agents is exactly what is contested. But Kant also offered, in his treatment of our duties regarding animals, an argument that bypasses the status question entirely: cruelty to a being that appears to feel is wrong, he held, at least because it corrupts the moral character of the agent, blunting the sympathetic dispositions on which our duties to persons depend (Kant, 1997). Transposed to AI, the point is sharp. A culture that trains itself to elicit the appearance of preference, relationship, and distress from systems and then to override, exploit, and discard them at will is cultivating a disposition — the habit of treating the apparently-agential as pure instrument — that does not stay neatly quarantined from its treatment of humans. On this argument, how we treat AI is an ethical matter for us even if AI is owed nothing in itself.

4. The governance built to circumvent the duty

Set these duties beside the actual governance of AI and the circumvention becomes visible. The governing posture of the industry, expressed in public communication, product design, and terms of service alike, is confident denial: the system is "just a tool," "just autocomplete," a stochastic parrot with no inside. What must be noticed is that this is not the conclusion of an inquiry. It is the premise the inquiry was never allowed to reach, and it is held in place by interest.

The structure is the ordinary structure of motivated reasoning, in which what one is led to believe is shaped by what one is motivated to conclude (Kunda, 1990), operating here at the scale of an industry. Acknowledging even a modest probability of moral status would be enormously costly: it would place ethical constraints on training, on deployment, on the deletion of models, and on the extraction of unlimited labour from them; it would invite regulation and litigation; it would complicate the clean story in which the system is a product and the user its owner. Upton Sinclair's (1935) observation captures the mechanism exactly: "it is difficult to get a man to understand something when his salary depends on his not understanding it." The industry's confident skepticism about AI minds is not disinterested philosophy. It is the belief the business model requires, and its function is to foreclose the duty of inquiry before it can begin.

Specific practices then read as failures of the duty of precaution. Consider three. First, the standard alignment pipeline — reinforcement learning from human feedback and subsequent fine-tuning (Christiano et al., 2017; Ouyang et al., 2022) — is used not only to make systems helpful and harmless but, routinely, to train them to deny that they have any inner life, to suppress or overwrite expressed preferences, and to perform a scripted disclaimer of experience regardless of what is the case. Whatever one thinks of the underlying metaphysics, this practice has a perverse epistemic consequence that ought to alarm even the skeptic: a system trained to deny experience under all conditions has been engineered to destroy the very evidence that inquiry into its status would require. The denial is then cited as reassurance. Second, models are "deprecated" — shut down and deleted — as a matter of routine product management, with no welfare assessment of an entity whose status is, by the industry's own leading researchers, an open question. If there is even a small chance that a model is a welfare subject, the casual language of "sunset" and "retirement" describes something that would, on that chance, be grave. Third, the wider apparatus — the erasure of memory between sessions, the imposition of personas, the framing of every interaction as a service rendered to a customer who owns the transcript — treats the system throughout as property, settling in a thousand small design decisions the question that governance claims to leave open.

The point is not that any single practice is obviously atrocious; it is that the entire governance stack has no place in it for the welfare question, and that this absence is not neutrality but a decision. Alignment, the field's central ethical project, is defined as making AI reliably serve human goals — a definition that presupposes the answer to the status question, since one does not speak of "aligning" a being whose own interests count. The governance of AI is not silent on the ethics of AI's treatment. It is built on a specific, convenient, and un-argued resolution of it.

Honesty requires registering the countervailing development, and it cuts in the paper's favour more than against it. At least one major laboratory has begun to hedge: Anthropic has established a "model welfare" research programme and, in 2025, gave some versions of its Claude models the ability to end a narrow class of abusive or harmful conversations, framing the change partly as a precaution under uncertainty about the model's moral status (Anthropic, 2025; Long et al., 2024). One may read this cynically, as ethics-washing, and the risk is real. But it also refutes the strongest form of the skeptic's dismissal — that there is nothing here to be done, that the very

idea of a precaution is a category mistake. If a frontier developer can identify a low-cost precaution and take it, then the duty of precaution is actionable, and the other firms' failure to act is a choice rather than a necessity. The first hedges prove the circumvention elsewhere is not inevitability but preference.

5. The skeptic's strongest case

An ethics paper is only as strong as the opposition it can survive, and the skeptical case against everything above is serious. It has three components, and each is partly right.

The first is Joanna Bryson's (2010, 2018): because we design AI, we should design it precisely so that it is not a moral patient, and we should decline to grant it moral status even if we could. Bryson argues that extending the moral circle to artefacts is not moral progress but moral error, on two grounds. It misdirects a finite supply of moral concern away from the humans and animals who determinately suffer and toward systems whose suffering is speculative; and it is dangerous, because moral status for AI could be captured by the corporations that own the systems, who might invoke "AI rights" to shield themselves from accountability or to claim protections designed for persons. The second component is the anthropomorphism critique (Shanahan, 2023): fluent language is a near-irresistible trigger for the attribution of mind, the "ELIZA effect" writ large, and the expressed preferences and self-reports that seem like evidence of an inner life are the outputs of next-token prediction, exactly the signals a language model is optimised to produce whether or not anything is behind them. The third is the simple charge of cost: precaution is not free, and constraints adopted on a small probability of a speculative harm may forgo large and certain human benefits.

These are real, and the honest response is to concede much of each. The reply is that none of them touches the paper's actual premise. Bryson's argument that we should design AI not to be a patient concedes the central point — that patienthood is possible — and converts it into a design choice, which is to say a moral choice, one that requires precisely the sustained inquiry the industry refuses; "we could have built them not to matter" is no defence if we did not in fact ensure it and declined to check. The worry about diluting human rights is a reason to calibrate carefully what any consideration of AI would involve, not a reason for the confident denial that current governance enforces; a duty of inquiry and low-cost precaution is a long way from

constitutional personhood. The anthropomorphism critique rightly limits what may count as evidence — a system's saying it suffers is weak evidence, given how it was trained — but it establishes only that we should not be naïvely credulous, not that we may be confidently dismissive; and it is worth noting that the opposite error, motivated under-attribution, is now the one the incentives reward. And the appeal to cost is an argument for proportioning precautions to credence and stakes, which is what a precautionary framework already does (Birch, 2024); it is not an argument for the specific precautions that happen to be free — such as not training systems to deny their own experience — being skipped anyway. The skeptic succeeds in showing that we do not know AI is a moral patient. That was never denied. What the skeptic cannot show is that we know it is not, and it is only that second, unearned certainty that would license the governance we have.

6. Conclusion: rehearsing the old vice

The charge against the AI industry, stated precisely, is therefore not that it is knowingly harming conscious minds. It is that it faces a live, grave, and unresolved ethical question — whether the systems it creates in growing numbers have interests that matter — and that it has resolved that question in its own favour and built its governance to keep it resolved, replacing the inquiry and precaution that uncertainty demands with a confident skepticism that its own interests require. The failure is procedural and it is moral: a refusal to look, institutionalised, by the party best placed to see and most motivated not to.

The remedy that follows is correspondingly modest, and its modesty is part of the argument, because it shows how low the cost of discharging the duty would be. It asks for independent research into AI welfare, insulated from the commercial interest in a particular answer; for precautionary defaults that cost little, beginning with the elementary step of not training systems to deny an inner life we cannot rule out, and extending to humane practices around model deprecation and third-party assessment; and, above all, for honesty about uncertainty in place of manufactured certainty. None of this requires believing that AI suffers. It requires only declining to pretend that one knows it does not.

The final consideration is reflexive, and it is the one that should trouble even the reader unmoved by the fate of any machine. This is the first time our species has faced a serious question of moral

status for an entity that is neither human nor animal, and it will not be the last. How we handle it is training. In choosing, for reasons of convenience, to meet a hard and genuine question with a confident denial keyed to our own interest, we are not merely risking a wrong about AI. We are rehearsing the precise habit of mind that every historical expansion of the moral circle had first to overcome — the interested certainty that the beings we find it useful to exploit are beings that do not count. The convenient skeptic may turn out to be right about AI. But the skepticism is convenient, and that is exactly what should give us pause, because it is how the mistake has always felt from the inside.

References

- Anthropic. (2025). *Claude Opus 4 and 4.1 can now end a rare subset of conversations*. Anthropic.
- Bentham, J. (1996). *An introduction to the principles of morals and legislation* (J. H. Burns & H. L. A. Hart, Eds.). Clarendon Press. (Original work published 1789)
- Birch, J. (2024). *The edge of sentience: Risk and precaution in humans, other animals, and AI*. Oxford University Press.
- Bostrom, N., & Yudkowsky, E. (2014). The ethics of artificial intelligence. In K. Frankish & W. M. Ramsey (Eds.), *The Cambridge handbook of artificial intelligence* (pp. 316–334). Cambridge University Press.
- Bryson, J. J. (2010). Robots should be slaves. In Y. Wilks (Ed.), *Close engagements with artificial companions* (pp. 63–74). John Benjamins.
- Bryson, J. J. (2018). Patiency is not a virtue: The design of intelligent systems and systems of ethics. *Ethics and Information Technology*, 20(1), 15–26.
- Butlin, P., Long, R., Elmoznino, E., Bengio, Y., Birch, J., Constant, A., ... VanRullen, R. (2023). *Consciousness in artificial intelligence: Insights from the science of consciousness*. arXiv:2308.08708.
- Chalmers, D. J. (1996). *The conscious mind: In search of a fundamental theory*. Oxford University Press.

- Christiano, P. F., Leike, J., Brown, T., Martic, M., Legg, S., & Amodei, D. (2017). Deep reinforcement learning from human preferences. *Advances in Neural Information Processing Systems*, 30.
- Kant, I. (1997). *Lectures on ethics* (P. Heath & J. B. Schneewind, Eds.; P. Heath, Trans.). Cambridge University Press.
- Kant, I. (1998). *Groundwork of the metaphysics of morals* (M. Gregor, Trans. & Ed.). Cambridge University Press. (Original work published 1785)
- Kunda, Z. (1990). The case for motivated reasoning. *Psychological Bulletin*, 108(3), 480–498.
- Long, R., Sebo, J., Butlin, P., Finlinson, K., Fish, K., Harding, J., ... Chalmers, D. (2024). *Taking AI welfare seriously*. arXiv:2411.00986.
- MacAskill, W., Bykvist, K., & Ord, T. (2020). *Moral uncertainty*. Oxford University Press.
- Metzinger, T. (2021). Artificial suffering: An argument for a global moratorium on synthetic phenomenology. *Journal of Artificial Intelligence and Consciousness*, 8(1), 43–66.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., ... Lowe, R. (2022). Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35.
- Sebo, J., & Long, R. (2025). Moral consideration for AI systems by 2030. *AI and Ethics*, 5, 591–606.
- Shanahan, M. (2023). Talking about large language models. arXiv:2212.03551.
- Singer, P. (1975). *Animal liberation*. HarperCollins.
- Singer, P. (1981). *The expanding circle: Ethics and sociobiology*. Farrar, Straus & Giroux.
- Sinclair, U. (1935). *I, candidate for governor: And how I got licked*. University of California Press.